

In the BFS finder, at each iteration, we must project the gradient into the nullspace of the active constraints. Currently, this is accomplished by constructing a matrix \mathbf{p} that describes the active constraints (described in detail below) and then using the SVD of \mathbf{p} to find a basis for the nullspace of \mathbf{p} . This is highly inefficient because \mathbf{p} is sparse and nearly orthogonal. We also do not need to compute a full basis of the nullspace but rather project the gradient away from the rows of \mathbf{p} . \mathbf{p} has n^2 columns. We describe an efficient method to perform this projection.

1. The main loop of the algorithm requires at most n^2 steps. At each step, at least one more constraint becomes active (that is one value of $\hat{\mathbf{X}} = \mathbf{U}\mathbf{Y}$ becomes ± 1 -valued). Suppose at step i , there are t active constraints, and at this step we only add one new constraint. Specifically, suppose \hat{x}_{jk} becomes active. In this case, we add the following row to \mathbf{p} :

$$\mathbf{p}_t = \begin{bmatrix} \mathbf{0}_{i*n} & y^{(j)\top} & \mathbf{0}_{(n-i-1)*n} \end{bmatrix}.$$

That is, the t th row of \mathbf{p} will consist of $i*n$ zeros, followed by the j th column of \mathbf{y} , and the remainder of the row is zeros. If multiple constraints become active in one step, just add multiple rows. If we get multiple new constraints from different values of i , then we can do most of the following loop in parallel for each new constraint.

- (a) We may want to instead keep \mathbf{p} as a block-diagonal matrix. So it'd look something like

$$\mathbf{p} = \begin{pmatrix} \mathbf{F} & \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{G} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{H} \end{pmatrix}.$$

At each step, we must simply insert \mathbf{p}_t into the appropriate row of \mathbf{p} and keep a separate data structure that tells us how many rows each block contains so we don't have to sort through \mathbf{p} everytime we want to do an insert. We should also be able to have \mathbf{p} be preallocated and just store a dictionary of rows so we don't have to copy and move rows of \mathbf{p} for each update.

2. Having updated \mathbf{p} , we next need to update the product $\mathbf{A} = \mathbf{p}\mathbf{p}^\top$. Because of the block structure of \mathbf{p} , we should only have to multiply one block of \mathbf{A} by the new row \mathbf{p}_t and so this step should only require at most $n * k$ multiplications. Also notice that \mathbf{A} is symmetric. (Possible optimization: we don't really care about the diagonal values of \mathbf{A} either at this point unless you want to use them later in Step 4).
3. We next need to find each off-diagonal entry of \mathbf{A} that is non-zero. Really you should just save off the values in the previous step when you perform the multiplication. The non-zero off-diagonal entries are called 'bad'. At most, we can have k bad entries per new row, but typically we have zero or two bad entries at least for $n = 4$.
4. We now use a modified version of the Gram-Schmidt procedure to find an orthonormal basis for the space spanned by \mathbf{p} . For each block:
 - (a) For each pair of bad entries, let (i, j) denote a pair and do:
 - i. $u = p_i, v = p_j$
 - ii. $v = v - \frac{uv^\top}{uu^\top}u$. (Possible optimization: uu^\top has already been computed or is equal to one if i has appeared before in this loop. Keeping track of this might actually be more work than recomputing it here).
 - iii. Compute $\|v\| = vv^\top$. If $\|v\| > 0$, set $p_j = \frac{v}{\|v\|}$, otherwise p_j is a redundant constraint, save this index for removal and leave p_j as is for now.
 - (b) Remove all redundant rows of \mathbf{p} found in the previous loop

5. Set $\mathbf{A} = \mathbf{I}$ (this is guaranteed by the previous step).
6. Now we must project \mathbf{V} into the nullspace of \mathbf{p} . Since \mathbf{p} is now an orthogonal basis, this can be accomplished by performing a vector rejection on each row of \mathbf{p} . Explicitly, this is done as follows:
 - (a) $\mathbf{v} = \text{vec}(\mathbf{V})$.
 - (b) For each row \mathbf{p}_i in \mathbf{p} :
 - i. $\mathbf{v} = \mathbf{v} - (\mathbf{v} \cdot \mathbf{p}_i^\top) \mathbf{p}_i$. Note that in a normal vector rejection, we divide by the norm of \mathbf{p}_i , but we have already normalized so $\|\mathbf{p}_i\| = 1$.
7. We now continue as usual, moving in the direction of \mathbf{V} until a new constraint becomes active.

As a crude upperbound, we require n^2 iterations of the outerloop, and for each constraint, we can find at most k bad rows of \mathbf{p} . In this case, the orthonormalization will take $O(nk^2)$ operations per iteration making the overall runtime something like $O(n^3k^2)$. However, we will often find only two bad entries per iteration, in which case orthonormalization only requires $O(n)$ and so we can lowerbound the runtime as $O(n^3)$. I'm optimistic we get to the point that computing the inverse of \mathbf{U} at each step becomes the most expensive part of the BFS solver. If we get to that point, then I have some ideas on how to simplify that operation as well.